

# Exam Try 2

March 13, 2018

## 1 Table of Contents

- 1 Generate  $P(j|i)$  role="presentation"> $P(j|i)P(j|i)$   
and find the histograms of edges for some node  $i$  role="presentation"> $i$ 
  - 2 Part 2, data
    - 2.1 Plot the data for three different perplexity values and describe.
    - 2.2 Perform NMF with a given  $K$ , plot a tSNE and color the data using the weightings. Then explore the weightings from NMF by finding the genes that contribute the most to each weighting.
    - 2.3 Take a look at the AML1 data

```
In [1]: import pandas as pd
import numpy as np
import scipy
import sklearn
from sklearn import manifold
from sklearn import decomposition

# Graphics
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import rc

rc('text', usetex=True)
rc('text.latex', preamble=r'\usepackage{cmbright}')
rc('font', **{'family': 'sans-serif', 'sans-serif': ['Helvetica']})

# Magic function to make matplotlib inline;
%matplotlib inline

# This enables SVG graphics inline.
# There is a bug, so uncomment if it works.
%config InlineBackend.figure_formats = {'png', 'retina'}

# JB's favorite Seaborn settings for notebooks
rc = {'lines.linewidth': 2,
      'axes.labelsize': 18,
```

```

        'axes.titlesize': 18,
        'axes.facecolor': 'DFDFE5'}
sns.set_context('notebook', rc=rc)
sns.set_style("dark")

mpl.rcParams['xtick.labelsize'] = 16
mpl.rcParams['ytick.labelsize'] = 16
mpl.rcParams['legend.fontsize'] = 14

```

## 2 Generate $P(j|i)$ and find the histograms of edges for some node $i$

I had initially started on problem 1, then found it to be full of over/underflow errors and tricky linear algebra, so now I opted for the bio track.

The function below implements the argument of the exponent,  $|y_i - y_j|^2 / \sigma_i$ . Then I just need to raise that to the power of Euler. Since the 'node' is unspecified, I generated some data, then found the node and plotted its histogram.

```

In [2]: def sne(X):
        """
        # calculate the dotproduct between each sample:
        # calculate  $|x_j|^2$  for each vector
        """
        sum_X = np.sum(np.square(X), 1)
        dotprod = -2 * np.dot(X, X.T)
        # calculate
        #  $|x_j|^2 - 2*x_i/x_j/\cos\theta = |x_j|^2 - 2*x_i - x_j|^2$ 
        # this is asymmetric
        Dprime = np.add(dotprod, sum_X)
        # symmetrize by completing the square:
        #  $|x_j|^2 - 2*x_i - x_j|^2 + |x_i|^2$ 
        D = np.add(Dprime.T, sum_X)

        # set  $D_{ii} = 0$ 
        D = D.astype(np.float)
        D = np.maximum(D, 0)
        np.fill_diagonal(D, 0)
        return D

In [3]: def shannon(data, sigma=1.0):
        """Given data (squared differences of vectors), return the entropy and  $p_{ij}$  values
        # Compute P-row and corresponding perplexity
        arg = -data/(2*sigma**2)

        if (arg > 0).any():
            raise ValueError('At least one probability is negative')

        if (arg > 710).any():
            raise ValueError('overflow warning, sigma={0:.2g}'.format(sigma))

```

```

P = np.exp(arg)
sumP = P.sum(axis=0)

#  $H = -\sum_j p_{ji} \log p_{ji}$ 
#  $p_{ji} = P / \text{sumP}$ 
#  $\log p_{ji} = \log P - \log \text{sumP}$ 
#  $H = \sum_j p_{ji} / \text{sumP} * (D_{ji} / 2 * \sigma^2 + \log(\text{sumP}))$ 
#  $H = \sum_j (p_{ji} * D_{ji} / 2 * \sigma^2) / \text{sumP} + p_{ji} / \text{sumP} * \log(\text{sumP})$ 
#  $H = \beta * \sum_j (p_{ji} * D_{ji}) / \text{sumP} + \sum_j p_{ji} / \text{sumP} * \log(\text{sumP})$ 
#  $\sum_j p_{ji} = \sum_j p(j|i) = 1$ 
#  $H = \beta * \text{meancondD} + \log(\text{sumP})$ 
H = np.log(sumP) + (2 * sigma ** 2) * np.sum(data * P) / sumP

if np.abs(H) == np.inf:
    raise ValueError('Entropy is undefined')

# normalize the  $p_{ij}$ 
P = P / sumP
return H, P

```

```

In [4]: def pca(X=np.array([]), no_dims=50):
        """PCA on X."""
        n, d = X.shape
        # mean center along columns
        X = X - X.mean(axis=0)
        # use eig for hermitian/symmetric matrices ;)
        L, W = np.linalg.eigh(np.dot(X.T, X))
        # transform the coordinates
        Y = np.dot(X, W[:, d-no_dims - 1:d - 1])
        return Y, L

```

```

In [5]: dims = 100
        g1 = np.random.normal(10, 2, dims)
        g2 = np.random.normal(9.99, 2, dims)
        g3 = np.random.normal(10, 3, dims)
        g4 = np.random.normal(10, 4, dims)

        mus = np.array([g1, g2, g3, g4])
        mus = np.abs(mus)

```

```

In [6]: # generate data points, complete with random errors:
        measurements = 500
        X = np.zeros((measurements, dims))
        cluster = np.zeros(measurements)

        for i in np.arange(measurements):
            choice = np.random.choice(len(mus), 1)[0]

```

```

sigma = np.abs(np.random.normal(0, choice, dims))
sample = np.random.normal(mus[choice], sigma)
sample[sample < 0] = 0
X[i, :] = sample
cluster[i] = choice

```

In [7]: *# normalize cells by counts*

```

X = X/X.sum(axis=1).reshape(X.shape[0], 1)*10**3
print(X.max(), X.sum(axis=1).max())

```

36.28033275967146 1000.0000000000005

In [8]: tX, l = pca(X)

In [9]: *def plot\_clusters(coords, jitter=0, \*\*kwargs):*

```

    for i in np.arange(4):
        x = coords[cluster == i, 0]
        y = coords[cluster == i, 1]
        if jitter:
            x += np.random.normal(0, jitter, len(x))
            y += np.random.normal(0, jitter, len(x))
        plt.scatter(x, y, label='Cluster {0}'.format(i), **kwargs)

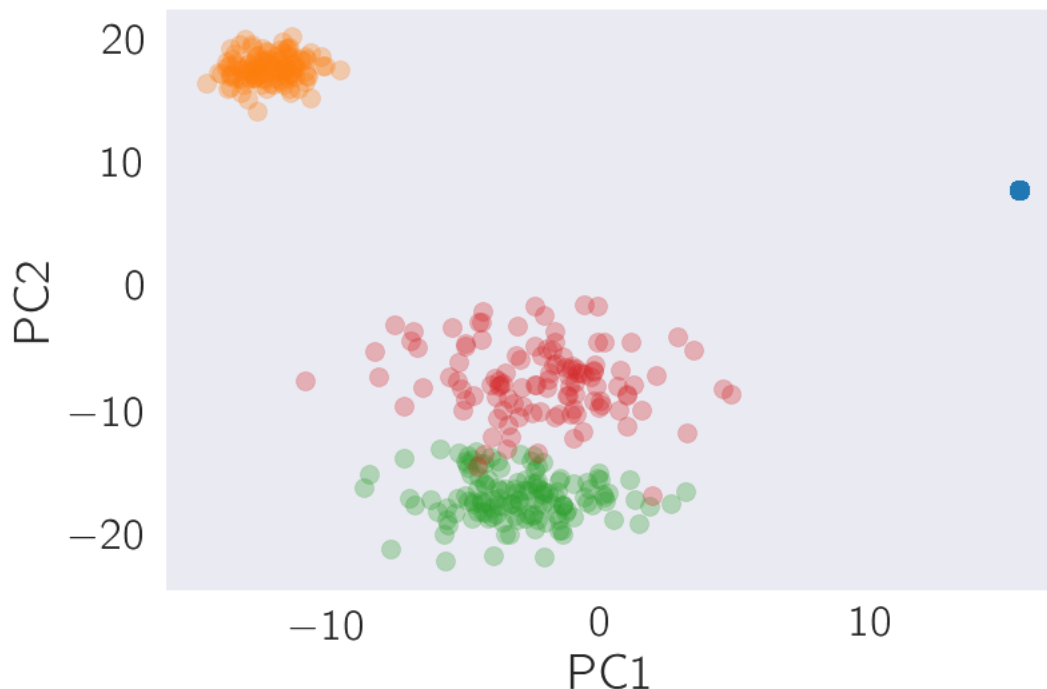
```

```

tX, l = pca(X, no_dims=2)
plot_clusters(tX, alpha=0.3)
plt.xlabel('PC1')
plt.ylabel('PC2')

```


Out [9]: Text(0,0.5,'PC2')



```
In [10]: H, P20 = shannon(sne(X)[0,:], sigma=20)
         H, P40 = shannon(sne(X)[0,:], sigma=40)
```

```
In [11]: sns.distplot(P20, label=r'\sigma=20$', kde=False)
         sns.distplot(P40, label=r'\sigma=40$', kde=False)
         plt.legend()
         plt.xlabel(r'$P(j|i)$')
```

```
Out[11]: Text(0.5,0,'$P(j|i)$')
```



### 3 Part 2, data

First, I will load my data. Then I will define a few useful functions. Then I will show all the required plots at once.

```
In [12]: names = pd.read_csv('../input/gene_names_class.csv', header=None)
```

```
In [13]: X = pd.read_csv('../input/healthy1.csv', header=None).as_matrix()
```

```
In [14]: def do_pca(x):
         normX = (x - x.mean(axis=1)[: , np.newaxis])/x.std(axis=1)[: , np.newaxis]
```

```

PCA = decomposition.PCA(n_components=10)
redX = PCA.fit_transform(normX.T)
return redX

def do_nmf(k, x):
    k = 5
    nnmfX = x/x.std(axis=1)[: , np.newaxis]
    model = decomposition.NMF(n_components=k)
    W = model.fit_transform(nnmfX)
    H = model.components_
    return W, H

def do_tsne(x, **kwargs):
    tsne = manifold.TSNE(n_components=2, **kwargs)
    Y = tsne.fit_transform(x)
    return Y

In [15]: def find_marker_genes(W):
    sig_genes = {}
    for i in np.arange(k):
        mu = W[:, i].mean()
        std = W[:, i].std()
        indices = np.where(W[:, i] > mu + std)
        index = np.where(W[:, i] == np.max(W[:, i]))[0][0]
        gene = names[names.index == index].values[0][0]
        sig_genes[i] = index
        print('{0} genes found for cluster {1}'.format(len(indices[0]), i))
        print('The gene with the best weight is {0}'.format(gene))
    return sig_genes

def plot_and_compare(k, tsne_X, normx, h, sig_genes):
    fig, ax = plt.subplots(ncols=k, nrows=2, figsize=(12, 6))
    fig.suptitle('t-SNE plot colored by gene expression of marker genes', fontsize=16)

    for i in np.arange(k):
        ax[0, i].scatter(tsne_X[:, 0], tsne_X[:, 1], c=normx[sig_genes[i]], alpha=0.3)
    ax[0, 0].set_ylabel('NMF weight')

    for i in np.arange(k):
        ax[1, i].scatter(tsne_X[:, 0], tsne_X[:, 1], c=h.T[:, i], alpha=0.3, s=15)
    ax[1, 0].set_ylabel('gene marker')

    plt.tight_layout()
    fig.tight_layout(rect=[0, 0.03, 1, 0.95])

```

### 3.1 Plot the data for three different perplexity values and describe.

```
In [17]: redX = do_pca(X)


fig, ax = plt.subplots(nrows=3, ncols=3, figsize=(9, 9))
perplexities = [2, 50, 100]
for i in np.arange(3):
    for j in np.arange(3):
        perp = perplexities[i]
        Y = do_tsne(redX, perplexity=perp)
        ax[j, i].scatter(Y[:, 0], Y[:, 1], s=10, alpha=0.3)

xlab = '$P=$' + '{0}'
ax[2, 0].set_xlabel(xlab.format(perplexities[0]))
ax[2, 1].set_xlabel(xlab.format(perplexities[1]))
ax[2, 2].set_xlabel(xlab.format(perplexities[2]))

fig.suptitle('3 different perplexity values', fontsize=20)

plt.tight_layout()
fig.tight_layout(rect=[0, 0.03, 1, 0.95])
```

3 different perplexity values



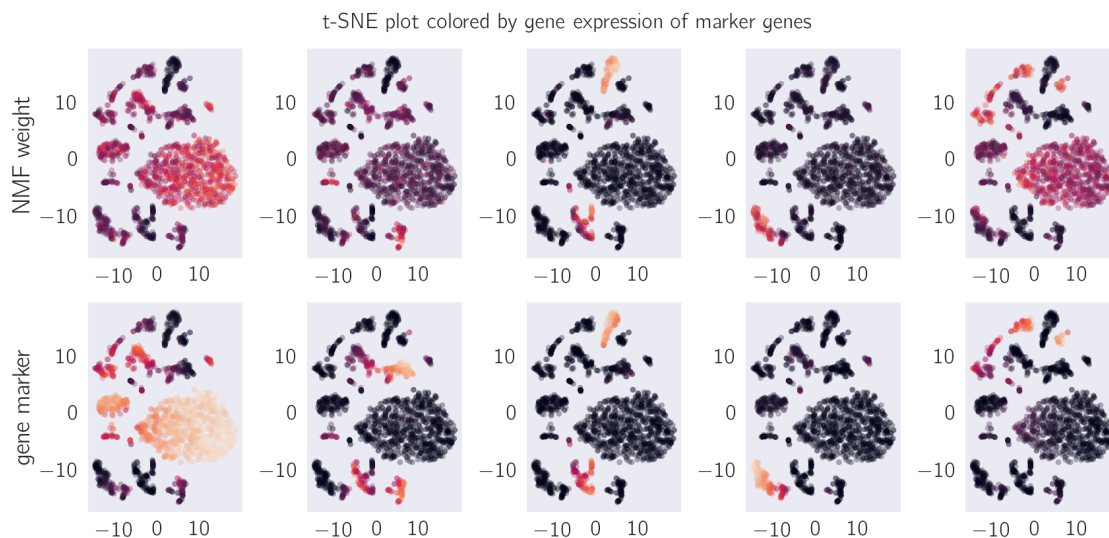
We can see that if the perplexity is too low, everything looks like a ball. For this dataset, the perplexity doesn't seem to matter very much between 50-100, since the diagrams look essentially identical to my eyes. In theory, setting the perplexity too high should make everything appear disjoint.

### 3.2 Perform NMF with a given K, plot a tSNE and color the data using the weightings. Then explore the weightings from NMF by finding the genes that contribute the most to each weighting.

```
In [18]: k = 5
         Y = do_tsne(redX)
         W, H = do_nmf(5, X)
         sig_genes = find_marker_genes(W)
         plot_and_compare(k, Y, X, H, sig_genes)
```



331 genes found for cluster 0  
 The gene with the best weight is RPS6  
 867 genes found for cluster 1  
 The gene with the best weight is PTMA  
 537 genes found for cluster 2  
 The gene with the best weight is HBB  
 565 genes found for cluster 3  
 The gene with the best weight is FTL  
 648 genes found for cluster 4  
 The gene with the best weight is B2M



Clearly, NMF and t-SNE both agree for this data. The clusters appear clear, though we ought to remember that distance means nothing in t-SNE space. Any separation at all between clusters means they might be totally distinct.

I found the genes that had the maximum weight within each cluster. Some of these genes are famous, and thus "make sense", which is a well known target of c-Myc, an immune protein...

Funnily enough, the largest cluster is characterized by expression of RPS6! Yay! What does this mean? Who knows. Nobody understands the ribosome.

### 3.3 Take a look at the AML1 data

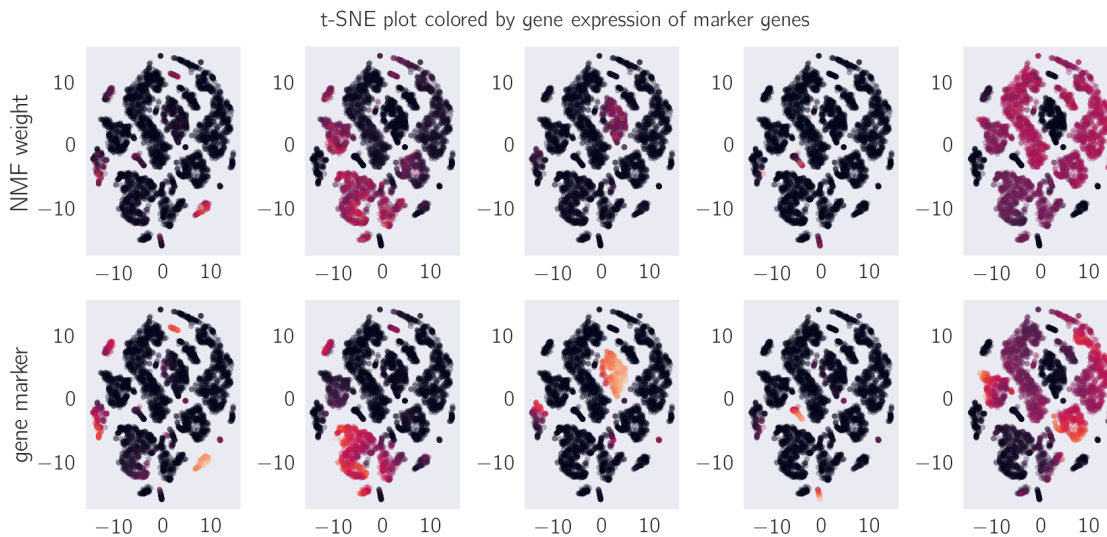
```
In [19]: unX = pd.read_csv('../input/aml1.csv', header=None).as_matrix()
          # clean up: there's some cells with nothing in here:
          unX = unX[np.where(unX.std(axis=1) != 0)[0], :]

In [20]: redUnX = do_pca(unX)
          unY = do_tsne(redUnX)
          W, unH = do_nmf(5, unX)
          sig_genes_un = find_marker_genes(W)
          plot_and_compare(k, unY, unX, unH, sig_genes_un)
```

```

848 genes found for cluster 0
The gene with the best weight is NPM1
762 genes found for cluster 1
The gene with the best weight is PYCARD
800 genes found for cluster 2
The gene with the best weight is BCAT2
741 genes found for cluster 3
The gene with the best weight is RPS6KB2
418 genes found for cluster 4
The gene with the best weight is NPRL3

```



Well, these beautiful images make no sense. Some clusters are observable, but the NMF colorings hardly make sense with these clusters. A major problem is that each color spans multiple clusters separated by uninterpretable space—does t-SNE think those are or are not clusters? Again, we get some "famous" genes that would make obvious sense to anyone who has studied them, and some ribosomal subunits that will make 99.9% of people shrug and go, 'yeah, we know that happens sometimes. When we get ribosomal subunits in genetic screens, we just kinda ignore them...'. Hey, maybe cancer cells have upregulated ribosomal subunits because they are actively growing! I'll make that my biological thought for the day.

In [ ]: